Speech-to-Text Models: Efficiency Snapshot (2026)

This table compares major speech-to-text (STT) models for analysing about 1 hour of audio. Efficiency is summarized using three aspects: approximate cost per hour (USD), relative accuracy (on general English, lower WER is better), and suitability for real-time subtitles/captions.

Provider	Model (2026)	~Cost for 1h audio	Relative accuracy*	Real-time captions?
OpenAI	Whisper API / GPT-4o Mini Transcribe	~$0.18–0.36	Very high on batch, multilingual	No (best for batch/offline)
AssemblyAI	Universal-3 Pro	~$0.15–0.21	Very high, among lowest WER on benchmarks	Yes (with U3 Pro Streaming)
Deepgram	Nova-3	~$0.26–0.31	Very high, state-of-the-art esp. streaming	Yes (sub-second latency)
Google Cloud	Speech-to-Text / Chirp	~$0.96	High, slightly behind STT specialists	Yes, but higher latency / more setup
Amazon	Transcribe Standard	~$1.44	High on clean audio, weaker in noisy tests	Yes, suitable for many AWS-native apps
Microsoft Azure	Azure Speech-to-Text	~$0.18–0.36	High, close to Google/AWS on English	Yes, with Azure Speech streaming

**Relative accuracy is based on public 2026 benchmarks and vendor data (WER on general English). Exact performance depends on your audio (language, noise, accents). Always benchmark on your own data.*